

# T-test one-way ANOVA\_ p-value concepts

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as ph
```

```
In [2]: from scipy import stats
import statistics as stat
```

```
In [3]: edu_p_school=[1200, 1400, 1600, 1800, 2000, 2200, 2400]
edu_h_school=[400, 500, 600, 650,700,750, 800, 900, 1000]
edu_p_college=[1100, 1300, 1500, 1700, 2900]
edu_Bechelore=[100, 200, 300, 400, 600, 1000, 1800]
edu_Graduate=[2500, 2600, 2700, 280, 2900, 3000, 3100]

print('p_school:' +str(int(stat.mean(edu_p_school))))
print('h_school:' +str(int(stat.mean(edu_h_school))))
print('p_college:' +str(int(stat.mean(edu_p_college))))
print('Bechelore:' +str(int(stat.mean(edu_Bechelore))))
print('p_Graduate:' +str(int(stat.mean(edu_Graduate))))

p_school:1800
h_school:700
p_college:1700
Bechelore:628
p_Graduate:2440
```

```
In [4]: t, p= stats.ttest_ind(edu_p_school, edu_p_college)

print(f't_value: {t}')
print(f'p_value: {p}')
```

```
t_value: 0.3057497440438928
p_value: 0.7660652423223867
```

```
In [5]: t1, p1= stats.ttest_ind(edu_h_school, edu_Bechelore)
```

```
print(f't_value: {t1}')
print(f'p_value: {p1}')
```

```
t_value: 0.3409594133082003
p_value: 0.7381969448005667
```


In [6]:

```
import seaborn as sns
sns.distplot(edu_Bachelore)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

Out[6]: <AxesSubplot:ylabel='Density'>



In [7]:

```
import seaborn as sns
sns.distplot(edu_Bachelore, label="Bachelore")
sns.distplot(edu_h_school, label="High school")
sns.distplot(edu_p_school, label="partial school")
sns.distplot(edu_p_college, label="Partial college")
sns.distplot(edu_Graduate, label="Graduate")
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

ated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
<AxesSubplot:ylabel='Density'>
```

Out[7]:



In [8]:

```
import seaborn as sns
import matplotlib.pyplot as plt
sns.distplot(edu_Bechelore, label="Bechelore")
sns.distplot(edu_h_school, label="High school")
sns.distplot(edu_p_school, label="partial school")
sns.distplot(edu_p_college, label="Partial college")
sns.distplot(edu_Graduate, label="Graduate")
plt.legend()
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

ction with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```

```
<matplotlib.legend.Legend at 0x1cac722ef10>
```

Out[8]:



In [9]:

```
f, p = stats.f_oneway(edu_Bechelore, edu_h_school, edu_p_school, edu_p_college, edu_Graduate)

print(f'f.value: {f}')
print(f'p.value: {p}')
```

```
f.value: 11.797908099455338
```

```
p.value: 7.0892689408058515e-06
```


```
In [10]: df=pd.read_csv("D:\\datasets\\insurance.csv")
df
```

```
Out[10]:
```


	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86
...	...	...	...	...	...	...	...
1333	50	male	31.0	3	no	northwest	10600.55
1334	18	female	31.9	0	no	northeast	2205.98
1335	18	female	36.9	0	no	southeast	1629.83
1336	21	female	25.8	0	no	southwest	2007.95
1337	61	female	29.1	0	yes	northwest	29141.36

1338 rows × 7 columns

```
In [11]: import seaborn as sns
sns.histplot(data=df, x="expenses", hue="sex", kde=True);
```



```
In [12]: sns.histplot(data=df, x="expenses", hue="smoker", kde=True);
```



```
In [13]: S_complet_y=df[df['smoker']=='yes']
          S_complet_n=df[df['smoker']=='no']

          stats.ttest_ind(S_complet_y['expenses'], S_complet_n['expenses'])
```

```
Out[13]: Ttest_indResult(statistic=46.6649210792002, pvalue=8.271449574495316e-283)
```

In [14]:

```
feature='smoker'
label='expenses'

groups=df[feature].unique()
for group in groups:
    print(group)
```

yes  
no

In [15]:

```
feature='smoker'
label='expenses'
groups=df[feature].unique()
grouped_values=[]
for group in groups:
    grouped_values.append(df[df[feature]==group][label])
```

grouped\_values

Out[15]:

```
[0      16884.92
 11      27808.73
 14      39611.76
 19      36837.47
 23      37701.88
 ...
1313     36397.58
1314     18765.88
1321     28101.33
1323     43896.38
1337     29141.36
Name: expenses, Length: 274, dtype: float64,
1       1725.55
2        4449.46
3       21984.47
4        3866.86
5        3756.62
 ...
1332     11411.69
1333     10600.55
1334      2205.98
1335      1629.83
```

```
1336      2007.95
Name: expenses, Length: 1064, dtype: float64]
```

```
In [16]: feature='smoker'
label='expenses'

groups=df[feature].unique()
grouped_values=[]
for group in groups:
    grouped_values.append(df[df[feature]==group][label])

grouped_values
stats.f_oneway(*grouped_values)
```

```
Out[16]: F_onewayResult(statistic=2177.6148593279827, pvalue=8.27144957450302e-283)
```

## *one-way ANOVA t-test bonferroni tukeyhsd\_ barplot*

```
In [17]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as ph

df=pd.read_csv("D:\\datasets\\bike_buyers.csv")
df.head()
```

```
Out[17]:
```

	ID	Marital_Status	Gender	Income	Children	Education	Occupation	Home_Owner	Cars	Commute_Distance	Region	Age	Purchased_Bi
0	12496	Married	Female	40000	1	Bachelors	Skilled Manual	Yes	0	0-1 Miles	Europe	42	
1	24107	Married	Male	30000	3	Partial College	Clerical	Yes	1	0-1 Miles	Europe	43	
2	14177	Married	Male	80000	5	Partial College	Professional	No	2	2-5 Miles	Europe	60	
3	24381	Single	Male	70000	0	Bachelors	Professional	Yes	1	5-10 Miles	Pacific	41	



	ID	Marital_Status	Gender	Income	Children	Education	Occupation	Home_Owner	Cars	Commute_Distance	Region	Age	Purchased_Bi
4	25597	Single	Male	30000	0	Bachelors	Clerical	No	0	0-1 Miles	Europe	36	,

In [18]: `df.Education.unique()`

Out[18]: `array(['Bachelors', 'Partial College', 'High School',  
'Partial High School', 'Graduate Degree'], dtype=object)`

In [19]: `df['Education_rank']=df['Education']  
df.Education_rank.replace(['Bachelors', 'Graduate Degree', 'High School', 'Partial College',  
'Partial High School'],[1, 2, 3, 4, 5], inplace=True)  
df.Education_rank.astype('int64')`

Out[19]:

0	1
1	4
2	4
3	1
4	1
	..
995	3
996	2
997	1
998	1
999	3

Name: Education\_rank, Length: 1000, dtype: int64

In [20]: `df['Commute_rank']=df['Commute_Distance']  
df.Commute_rank.replace(['0-1 Miles', '1-2 Miles', '2-5 Miles', '5-10 Miles', '10+ Miles'],[0, 1, 2, 5, 10], inplace=True)  
df.Commute_rank.astype('int64')`

Out[20]:

0	0
1	0
2	2
3	5
4	0
	..
995	2
996	2
997	0

```

998     1
999    10
Name: Commute_rank, Length: 1000, dtype: int64

```

```

In [21]: def clean_bike_buyers():
import pandas as pd
df=pd.read_csv("D:\\datasets\\bike_buyers.csv")
df['Education_rank']=df['Education']
df['Commute_rank']=df['Commute_Distance']
df['Purchased_Bike']=df['Purchased_Bike']

df.Education_rank.replace(['Bachelors', 'Graduate Degree', 'High School', 'Partial College',
                           'Partial High School'],[1, 2, 3, 4, 5], inplace=True)
df.Commute_rank.replace(['0-1 Miles', '1-2 Miles', '2-5 Miles', '5-10 Miles', '10+ Miles'],[0, 1, 2, 5, 10], inplace=True)
df['Purchased_Bike'].replace(['Yes', 'No'], [0,1], inplace=True)

df.astype({'Education_rank': 'int64'})
df.astype({'Commute_rank': 'int64'})
df['Purchased_Bike'].astype('int64')
return df

```

```

In [22]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as ph

df= clean_bike_buyers()

```

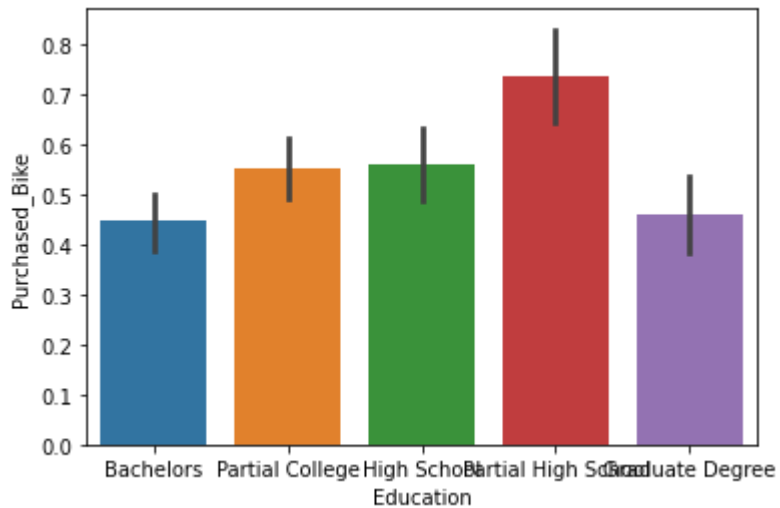
```

In [23]: import seaborn as sns
sns.barplot(df['Education'], df['Purchased_Bike']);


```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```




```
In [24]: plt.figure(figsize=(10, 5))
sns.barplot(x=df['Education'], y=df['Purchased_Bike']);
```




```
In [25]: viz = sns.barplot(x=df['Education'], y=df['Purchased_Bike'])
viz.set_xticklabels(viz.get_xticklabels(), rotation=25)
```


```
Out[25]: [Text(0, 0, 'Bachelors'),  
Text(1, 0, 'Partial College'),  
Text(2, 0, 'High School'),  
Text(3, 0, 'Partial High School'),  
Text(4, 0, 'Graduate Degree')]
```



```
In [26]: viz =sns.barplot(x=df['Education'], y=df['Purchased_Bike'],  
                        order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);  
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```




```
In [27]: viz = sns.barplot(x=df['Education'], y=df['Purchased_Bike'], estimator=np.median,  
                        order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);  
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```




In [28]:

```
viz =sns.barplot(x=df['Education'], y=df['Purchased_Bike'], estimator=np.median, ci='sd',  
                order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);  
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```

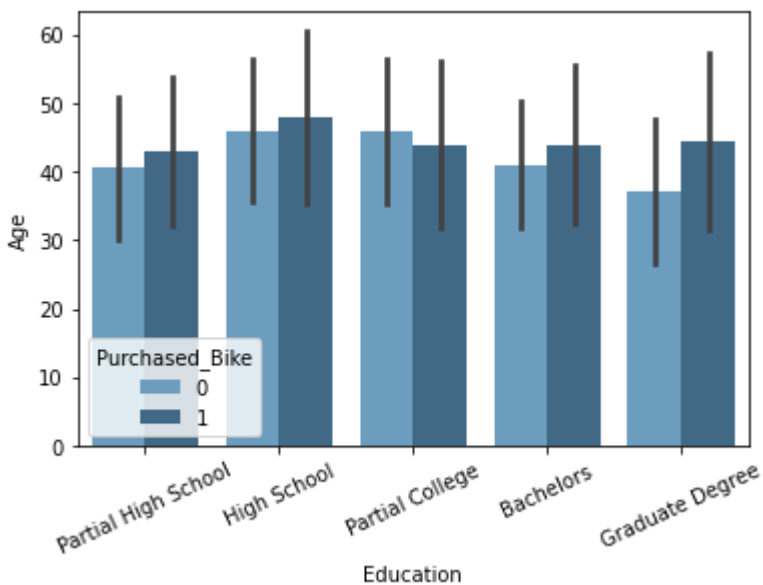


In [29]:


```
viz =sns.barplot(x=df['Education'], y=df['Purchased_Bike'], estimator=np.median, ci='sd', palette='Blues_d',  
                order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);  
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```



```
In [30]: viz = sns.barplot(x=df['Education'], y=df['Age'], hue=df['Purchased_Bike'], estimator=np.median, ci='sd', palette='Blues_c',
                        order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```




```
In [31]: viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', color='Home_Owner',
                        estimator=np.median, ci='sd', palette='Blues_d', order=['Partial High School', 'High School',
                                                                              'Partial College', 'Bachelors', 'Graduate Degree'],
                        viz.set_xticklabels(rotation=25);
```




```
In [32]: viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', col='Home_Owner',
                        estimator=np.median, ci='sd', order=['Partial High School', 'High School',
                                                                              'Partial College', 'Bachelors', 'Graduate Degree']));
                        viz.set_xticklabels(rotation=25);
```






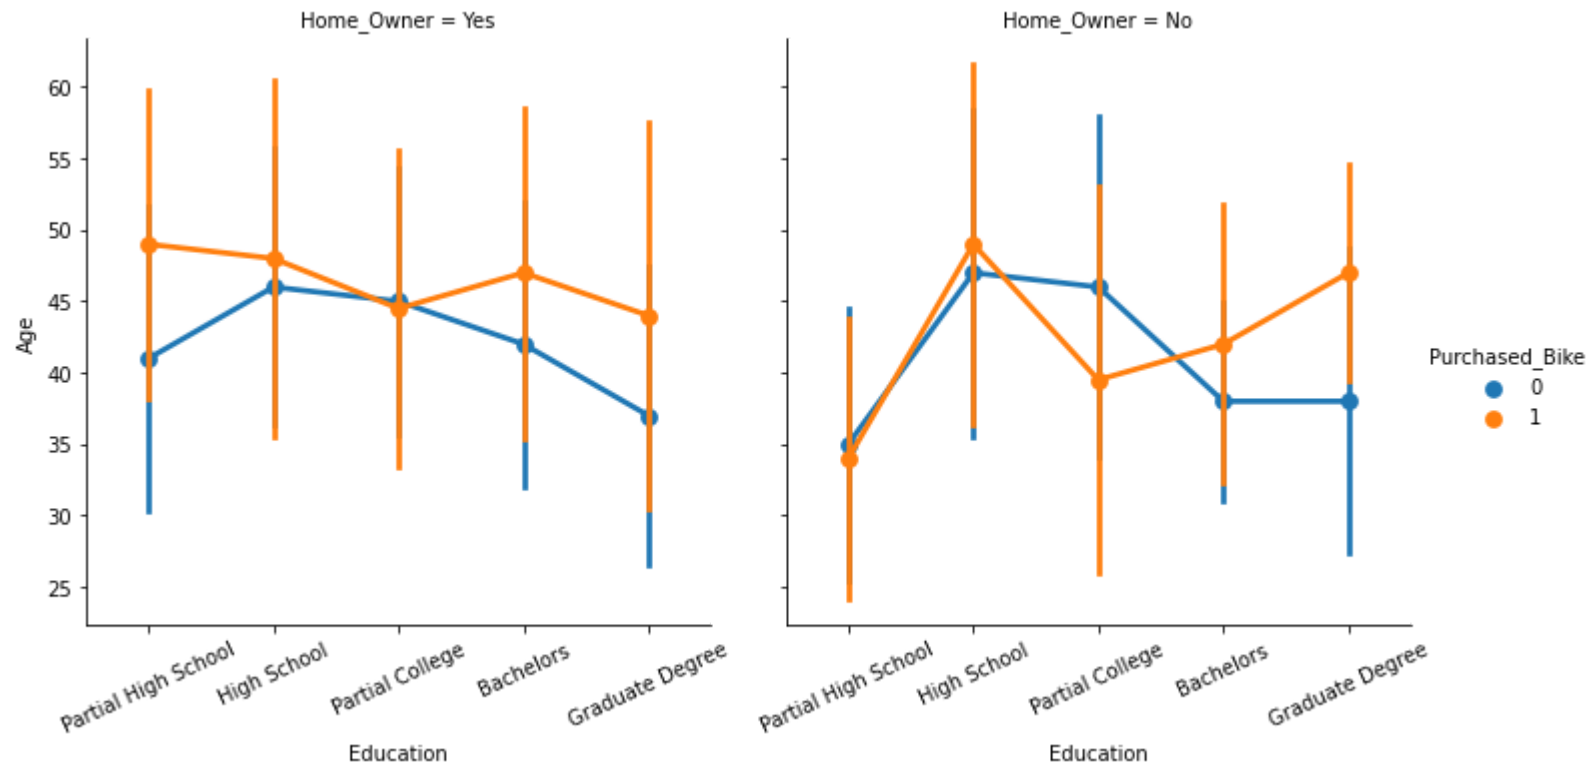
```
In [33]: viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', col='Home_Owner',
                        estimator=np.median, ci='sd', kind='boxen', order=['Partial High School', 'High School',
                        'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(rotation=25);
```



```
In [34]: viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', col='Home_Owner',
                        estimator=np.median, ci='sd', kind='box', order=['Partial High School', 'High School',
                        'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(rotation=25);
```



```
In [35]: viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', col='Home_Owner',
                        estimator=np.median, ci='sd', kind='point', order=['Partial High School', 'High School',
                        'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(rotation=25);
```



In [36]:

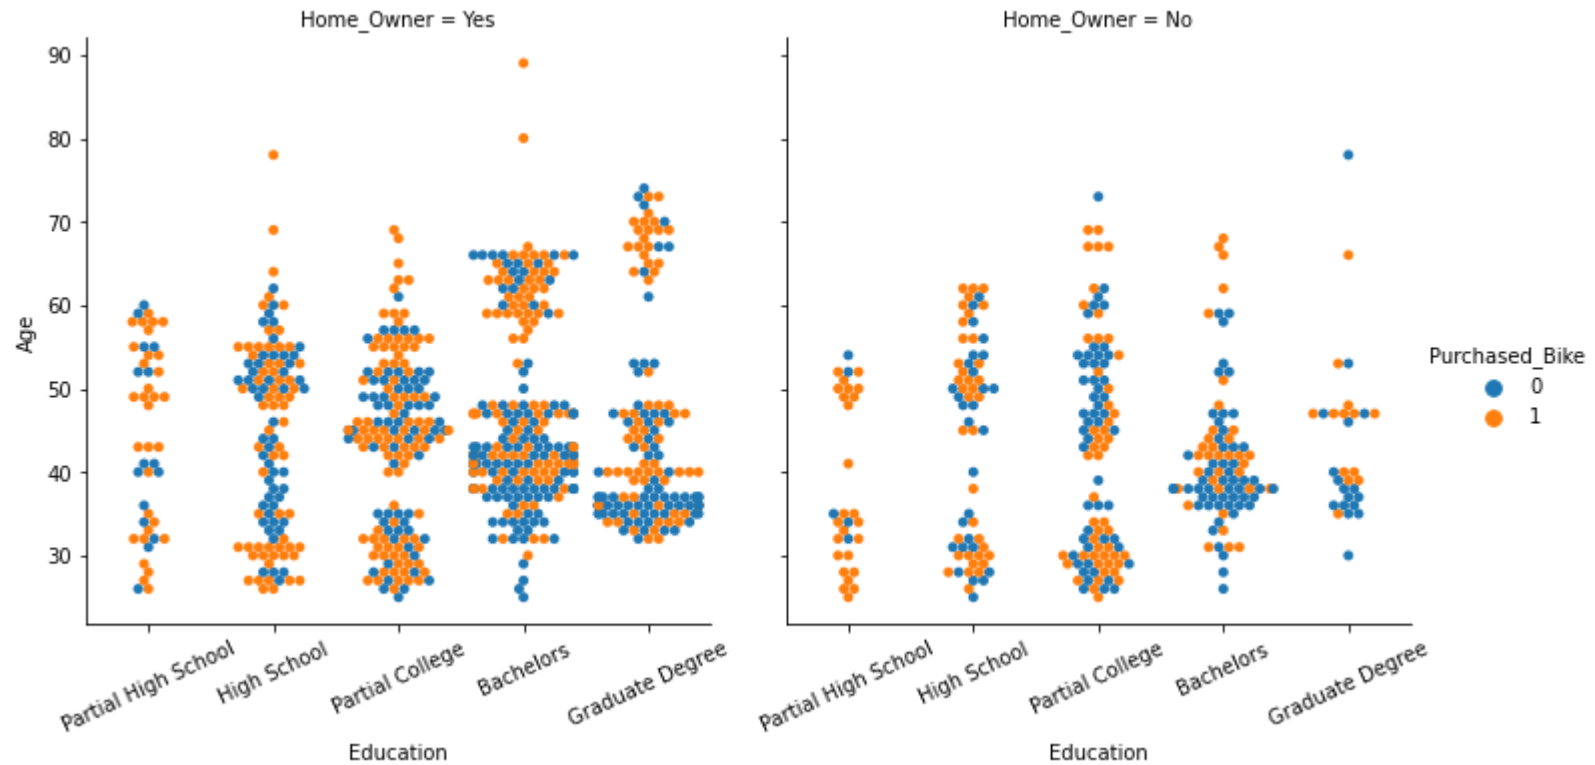
```
viz = sns.catplot(data=df, x='Education', y='Age', hue='Purchased_Bike', col='Home_Owner', estimator=np.median,
ci='sd', kind='swarm', order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(rotation=25);
```

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 16.1% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)

C:\Users\TAWAB COMPUTERS\anaconda3\lib\site-packages\seaborn\categorical.py:1296: UserWarning: 12.5% of the points cannot be placed; you may want to decrease the size of the markers or use stripplot.

warnings.warn(msg, UserWarning)



In [37]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

df= pd.DataFrame({'Locations':['Atlanta', 'Los Angeles', 'New York City', 'Phoenix'],
                  'Q1 Sales':[1567811, 3391023, 3409871, 789123],
                  'Q2 Sales':[1981237, 3609877, 3100098, 810988],
                  'Q3 Sales':[1761231, 3509889, 3209876, 751233],
                  'Q4 Sales':[3578500, 6712333, 3378900, 1500092]}))

#creat the position of the bar
x= np.arange(len(df.Locations))
# store the three columns from the DataFrame and "flatten" them
#to appear as regular python list structure
list_1 = df['Q1 Sales'].values.flatten()
list_2 = df['Q2 Sales'].values.flatten()
list_3 = df['Q3 Sales'].values.flatten()
list_4 = df['Q4 Sales'].values.flatten()
```

```

#plot the pokem names as the x ticks
plt.xticks(x, df.Locations)

# creat a Legend
plt.legend(loc='upper right')

#Add Label and title
plt.xlabel('Markets')
plt.ylabel('Sales in Millions')
plt.title('Sales by Quarter and Location')

#add an sns style and increase figure size
sns.set_style('white')
sns.set_context({"figure.figsize": (4, 6)})
sns.despine(top=True, right=True)


#show the plot
plt.show()
#Plot thr Barchat
plt.bar(x, list_1, label='Q1')
plt.bar(x, list_2, bottom=list_1, label='Q2')
plt.bar(x, list_3, bottom=list_1+list_2, label='Q3')
plt.bar(x, list_4, bottom=list_1+list_2+list_3, label='Q4')

```

No handles with labels found to put in legend.



Out[37]: <BarContainer object of 4 artists>



In [38]:

```

from scipy import stats
import pandas as pd

df= clean_bike_buyers()

groups= df['Education'].unique()
group_labels=[]
for g in groups:
    group_labels.append(df[df["Education"]==g]["Purchased_Bike"])

#now calcule the ANOVA results
F, p =stats.f_oneway(*group_labels)

print('F: ' +str(round(F, 4)))
print('p: ' +str(round(p, 4)))

```

F: 6.4653

p: 0.0


```
In [39]: r, p = stats.pearsonr(df["Education_rank"], df["Purchased_Bike"])

print('r: ' + str(round(r, 4)))
print('p: ' + str(round(p, 4)))
```

r: 0.1413

p: 0.0

```
In [40]: viz =sns.barplot(x=df['Education'], y=df['Purchased_Bike'],
                        order=['Partial High School', 'High School', 'Partial College', 'Bachelors', 'Graduate Degree']);
viz.set_xticklabels(viz.get_xticklabels(), rotation=25);
```



```
In [41]: Partial_high_school=df[df.Education== 'Partial High School']
High_school=df[df.Education== 'High School']
t, p= stats.ttest_ind(Partial_high_school['Purchased_Bike'], High_school['Purchased_Bike'])
```



```
print('t: ' + str(round(t, 2)))
print('p: ' + str(round(p, 2)))
```

t: 2.7

p: 0.01

In [42]: `from statsmodels.stats.multicomp import MultiComparison`

```
mc= MultiComparison(df['Purchased_Bike'], df['Education'])
print(mc.tukeyhsd())
```

```
Multiple Comparison of Means - Tukey HSD, FWER=0.05
=====
```

group1	group2	meandiff	p-adj	lower	upper	reject
Bachelors	Graduate Degree	0.0121	0.9	-0.1163	0.1404	False
Bachelors	High School	0.1109	0.1206	-0.0162	0.2381	False
Bachelors	Partial College	0.1032	0.0941	-0.0102	0.2166	False
Bachelors	Partial High School	0.2891	0.001	0.1159	0.4623	True
Graduate Degree	High School	0.0989	0.3296	-0.045	0.2428	False
Graduate Degree	Partial College	0.0912	0.3235	-0.0407	0.223	False
Graduate Degree	Partial High School	0.2771	0.001	0.0913	0.4629	True
High School	Partial College	-0.0077	0.9	-0.1385	0.123	False
High School	Partial High School	0.1782	0.0656	-0.0068	0.3632	False
Partial College	Partial High School	0.1859	0.0322	0.0101	0.3617	True

```
-----
```

In [43]: `e_types= df.Education.unique()`  
`ttests= []`

```
for i, e in enumerate(e_types):
    print(i, '', e)
```

```
0  Bachelors
1  Partial College
2  High School
3  Partial High School
4  Graduate Degree
```

In [44]: `e_types= df.Education.unique()`  
`ttests= []`

```
for i, e in enumerate(e_types):
    for i2, e2 in enumerate(e_types):
```

```

if i2>i:
    g1 = df[df.Education==e]['Purchased_Bike']
    g2 = df[df.Education==e2]['Purchased_Bike']
    t, p = stats.ttest_ind(g1, g2)
    ttests.append(f'{e} - {e2}: {t.round(4)}, {p.round(4)}')
    print(f'{e} - {e2}: {t.round(4)}, {p.round(4)}')
#print(i, '', '', e)

```

```

Bachelors - Partial College: -2.4693, 0.0138
Bachelors - High School: -2.3675, 0.0183
Bachelors - Partial High School: -4.6253, 0.0
Bachelors - Graduate Degree: -0.2546, 0.7991
Partial College - High School: -0.1601, 0.8729
Partial College - Partial High School: -2.9354, 0.0036
Partial College - Graduate Degree: 1.8728, 0.0618
High School - Partial High School: -2.698, 0.0074
High School - Graduate Degree: 1.862, 0.0634
Partial High School - Graduate Degree: 4.1685, 0.0

```

In [45]:

```

e_types= df.Education.unique()
ttests= []

for i, e in enumerate(e_types):
    for i2, e2 in enumerate(e_types):
        if i2>i:
            g1 = df[df.Education==e]['Purchased_Bike']
            g2 = df[df.Education==e2]['Purchased_Bike']
            t, p = stats.ttest_ind(g1, g2)
            ttests.append(f'{e} - {e2}: {t.round(4)}, {p.round(4)}')

threshold = 0.05/ len(ttests)
print(f'Significant t-test below {threshold}:')
for t in ttests:
    if t[2] <= threshold:
        print(t)

```

Significant t-test below 0.005:

```

-----
TypeError                                Traceback (most recent call last)
C:\Users\TAWABC~1\AppData\Local\Temp\ipykernel_2528\2544714263.py in <module>
    13 print(f'Significant t-test below {threshold}:')
    14 for t in ttests:
--> 15     if t[2] <= threshold:
    16         print(t)

```

**TypeError:** '<=' not supported between instances of 'str' and 'float'

In [ ]:

## Intro to MLR OLS in statmodels.api

In [46]:

```
import pandas as pd, numpy as np, statsmodels.api as sm
```

In [47]:

```
df=pd.read_csv("D:\\datasets\\insurance.csv")
df.head()
```

Out[47]:

	age	sex	bmi	children	smoker	region	expenses
0	19	female	27.9	0	yes	southwest	16884.92
1	18	male	33.8	1	no	southeast	1725.55
2	28	male	33.0	3	no	southeast	4449.46
3	33	male	22.7	0	no	northwest	21984.47
4	32	male	28.9	0	no	northwest	3866.86

In [48]:

```
label = "expenses"

y=df.expenses
x=df[['age', 'bmi', 'children']].assign(const=1)

model= sm.OLS(y, x)
results= model.fit()
print(results.summary())
```

### OLS Regression Results

```
=====
Dep. Variable:          expenses    R-squared:            0.120
Model:                  OLS        Adj. R-squared:       0.118
Method:                 Least Squares    F-statistic:         60.74
Date:                  Thu, 10 Aug 2023    Prob (F-statistic):   8.32e-37
```

```

Time:                16:21:39    Log-Likelihood:    -14392.
No. Observations:    1338        AIC:                2.879e+04
Df Residuals:        1334        BIC:                2.881e+04
Df Model:            3
Covariance Type:      nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
age          239.9626      22.288      10.766      0.000      196.239      283.686
bmi          332.5216      51.307       6.481      0.000      231.870      433.173
children     543.0436     258.230       2.103      0.036       36.462     1049.625
const     -6929.3145     1757.434      -3.943      0.000     -1.04e+04     -3481.678
=====
Omnibus:            325.223    Durbin-Watson:           2.012
Prob(Omnibus):      0.000    Jarque-Bera (JB):        602.850
Skew:               1.520    Prob(JB):               1.24e-131
Kurtosis:           4.254    Cond. No.                290.
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

In [49]: df['predictions']=results.fittedvalues
df

```

```

Out[49]:

```

	age	sex	bmi	children	smoker	region	expenses	predictions
0	19	female	27.9	0	yes	southwest	16884.92	6907.326136
1	18	male	33.8	1	no	southeast	1725.55	9172.284502
2	28	male	33.0	3	no	southeast	4449.46	12391.979997
3	33	male	22.7	0	no	northwest	21984.47	8537.689692
4	32	male	28.9	0	no	northwest	3866.86	10359.360926
...	...	...	...	...	...	...	...	...
1333	50	male	31.0	3	no	northwest	10600.55	17006.113044
1334	18	female	31.9	0	no	northeast	2205.98	7997.449896
1335	18	female	36.9	0	no	southeast	1629.83	9660.057790
1336	21	female	25.8	0	no	southwest	2007.95	6688.955930

	age	sex	bmi	children	smoker	region	expenses	predictions
1337	61	female	29.1	0	yes	northwest	29141.36	17384.779329

1338 rows × 8 columns

```
In [50]: print(results.predict([19, 27.9, 0, 1]))
```

[6907.32613573]

## MLR with categorical values dummy codes

```
In [51]: for col in df:
         if not pd.api.types.is_numeric_dtype(df[col]):
             df = pd.get_dummies(df, columns=[col], drop_first=True)
         df.head()
```

```
Out[51]:
```

	age	bmi	children	expenses	predictions	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
0	19	27.9	0	16884.92	6907.326136	0	1	0	0	1
1	18	33.8	1	1725.55	9172.284502	1	0	0	1	0
2	28	33.0	3	4449.46	12391.979997	1	0	0	1	0
3	33	22.7	0	21984.47	8537.689692	1	0	1	0	0
4	32	28.9	0	3866.86	10359.360926	1	0	1	0	0

```
In [52]: import pandas as pd, numpy as np, statsmodels.api as sm
         x = df.drop(columns=[label]).assign(const=1)
         results = sm.OLS(y, x).fit()
         print(results.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          expenses    R-squared:                0.751
Model:                  OLS        Adj. R-squared:             0.749
Method:                 Least Squares    F-statistic:            500.9
Date:                  Thu, 10 Aug 2023    Prob (F-statistic):      0.00

```

```

Time:                16:21:40    Log-Likelihood:        -13548.
No. Observations:    1338        AIC:                2.711e+04
Df Residuals:        1329        BIC:                2.716e+04
Df Model:            8
Covariance Type:     nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
age           -154.8900      31.521      -4.914      0.000     -216.727     -93.053
bmi           -231.2527      30.546      -7.571      0.000     -291.177    -171.328
children      -456.0686     149.330      -3.054      0.002     -749.018    -163.120
predictions         1.7158       0.142     12.122      0.000         1.438         1.993
sex_male      -131.3520     332.935      -0.395      0.693     -784.488     521.784
smoker_yes     2.385e+04     413.139     57.723      0.000      2.3e+04     2.47e+04
region_northwest -352.7901     476.261      -0.741      0.459    -1287.095     581.515
region_southeast -1035.5957     478.681      -2.163      0.031    -1974.648     -96.544
region_southwest -959.3058     477.912      -2.007      0.045    -1896.850     -21.762
const          -52.2030      12.612      -4.139      0.000      -76.944     -27.462

```

```

=====
Omnibus:                300.499    Durbin-Watson:           2.088
Prob(Omnibus):          0.000    Jarque-Bera (JB):        719.382
Skew:                   1.212    Prob(JB):                6.14e-157
Kurtosis:               5.652    Cond. No.                2.34e+19
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.  
[2] The smallest eigenvalue is 4.75e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

## MLR OLS *standardization* normalization

In [53]:

```

from sklearn import preprocessing

df_zscore= pd.DataFrame(preprocessing.StandardScaler().fit_transform(df), columns=df.columns)
df_zscore.head()

```

Out[53]:

	age	bmi	children	expenses	predictions	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
0	-1.438764	-0.453646	-0.908614	0.298583	-1.516295	-1.010519	1.970587	-0.566418	-0.611324	1.765481
1	-1.509965	0.514186	-0.078767	-0.953689	-0.976566	0.989591	-0.507463	-0.566418	1.635795	-0.566418

	age	bmi	children	expenses	predictions	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
2	-0.797954	0.382954	1.580926	-0.728675	-0.209329	0.989591	-0.507463	-0.566418	1.635795	-0.566418
3	-0.441948	-1.306650	-0.908614	0.719843	-1.127787	0.989591	-0.507463	1.765481	-0.611324	-0.566418
4	-0.513149	-0.289606	-0.908614	-0.776802	-0.693692	0.989591	-0.507463	1.765481	-0.611324	-0.566418

In [54]:

```

y= df_zscore.expenses

x=df_zscore.drop(columns=['predictions', 'expenses']).assign(const=1)

model= sm.OLS(y, x)
results= model.fit()
print(results.summary())

```

## OLS Regression Results

```

=====
Dep. Variable:          expenses    R-squared:                0.751
Model:                  OLS        Adj. R-squared:           0.749
Method:                 Least Squares    F-statistic:             500.9
Date:                  Thu, 10 Aug 2023    Prob (F-statistic):       0.00
Time:                  16:21:41          Log-Likelihood:          -968.62
No. Observations:      1338             AIC:                   1955.
Df Residuals:          1329             BIC:                   2002.
Df Model:               8
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
age	0.2980	0.014	21.586	0.000	0.271	0.325
bmi	0.1709	0.014	11.864	0.000	0.143	0.199
children	0.0474	0.014	3.452	0.001	0.020	0.074
sex_male	-0.0054	0.014	-0.395	0.693	-0.032	0.022
smoker_yes	0.7950	0.014	57.723	0.000	0.768	0.822
region_northwest	-0.0125	0.017	-0.741	0.459	-0.046	0.021
region_southeast	-0.0381	0.018	-2.163	0.031	-0.073	-0.004
region_southwest	-0.0340	0.017	-2.007	0.045	-0.067	-0.001
const	3.296e-17	0.014	2.41e-15	1.000	-0.027	0.027

```

=====
Omnibus:                 300.499    Durbin-Watson:           2.088
Prob(Omnibus):           0.000     Jarque-Bera (JB):        719.382
Skew:                    1.212     Prob(JB):                6.14e-157
Kurtosis:                 5.652     Cond. No.                 2.21
=====


```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [55]:


```
import seaborn as sns  
  
sns.histplot(y);
```



In [56]:

```
sns.histplot(df.expenses);
```





```
In [57]: df_minmax= pd.DataFrame(preprocessing.MinMaxScaler().fit_transform(df), columns=df.columns)
df_minmax.head()
```

```
Out[57]:
```

	age	bmi	children	expenses	predictions	sex_male	smoker_yes	region_northwest	region_southeast	region_southwest
0	0.021739	0.320755	0.0	0.251611	0.198761	0.0	1.0	0.0	0.0	1.0
1	0.000000	0.479784	0.2	0.009636	0.306026	1.0	0.0	0.0	1.0	0.0
2	0.217391	0.458221	0.6	0.053115	0.458506	1.0	0.0	0.0	1.0	0.0
3	0.326087	0.180593	0.0	0.333010	0.275973	1.0	0.0	1.0	0.0	0.0
4	0.304348	0.347709	0.0	0.043816	0.362244	1.0	0.0	1.0	0.0	0.0

```
In [58]: y= df_minmax.expenses

x=df_minmax.drop(columns=['predictions', 'expenses']).assign(const=1)

model= sm.OLS(y, x)
```

```
results= model.fit()
print(results.summary())
```

```

=====
                        OLS Regression Results
=====
Dep. Variable:          expenses      R-squared:            0.751
Model:                  OLS          Adj. R-squared:       0.749
Method:                 Least Squares  F-statistic:         500.9
Date:                   Thu, 10 Aug 2023  Prob (F-statistic):    0.00
Time:                   16:21:41      Log-Likelihood:      1230.9
No. Observations:      1338          AIC:                 -2444.
Df Residuals:          1329          BIC:                 -2397.
Df Model:               8
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
age	0.1886	0.009	21.586	0.000	0.171	0.206
bmi	0.2009	0.017	11.864	0.000	0.168	0.234
children	0.0380	0.011	3.452	0.001	0.016	0.060
sex_male	-0.0021	0.005	-0.395	0.693	-0.013	0.008
smoker_yes	0.3807	0.007	57.723	0.000	0.368	0.394
region_northwest	-0.0056	0.008	-0.741	0.459	-0.021	0.009
region_southeast	-0.0165	0.008	-2.163	0.031	-0.032	-0.002
region_southwest	-0.0153	0.008	-2.007	0.045	-0.030	-0.000
const	-0.0481	0.009	-5.137	0.000	-0.066	-0.030

```

=====
Omnibus:                 300.499      Durbin-Watson:          2.088
Prob(Omnibus):           0.000      Jarque-Bera (JB):       719.382
Skew:                    1.212      Prob(JB):               6.14e-157
Kurtosis:                 5.652      Cond. No.                9.58
=====

```


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


## \_MLR\_OLS assumptions normality multicollinearity VIF

```
In [59]: y=np.log1p(y)
sns.histplot(y)
```

```
Out[59]: <AxesSubplot:xlabel='expenses', ylabel='Count'>
```



```
In [60]: sns.histplot(np.log(df.expenses));
```



```
In [61]: y= np.log(df.expenses)
x= df.drop(columns=['predictions', 'expenses']).assign(const=1)

print(sm.OLS(y, x).fit().summary())
```

#### OLS Regression Results

```
=====
Dep. Variable:          expenses    R-squared:                0.768
Model:                  OLS        Adj. R-squared:           0.767
Method:                 Least Squares    F-statistic:             549.7
Date:                   Thu, 10 Aug 2023    Prob (F-statistic):      0.00
Time:                   16:21:42          Log-Likelihood:          -808.54
No. Observations:      1338              AIC:                    1635.
Df Residuals:          1329              BIC:                    1682.
Df Model:               8
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----						

age	0.0346	0.001	39.654	0.000	0.033	0.036
bmi	0.0134	0.002	6.377	0.000	0.009	0.017
children	0.1019	0.010	10.086	0.000	0.082	0.122
sex_male	-0.0754	0.024	-3.090	0.002	-0.123	-0.028
smoker_yes	1.5543	0.030	51.330	0.000	1.495	1.614
region_northwest	-0.0638	0.035	-1.827	0.068	-0.132	0.005
region_southeast	-0.1572	0.035	-4.480	0.000	-0.226	-0.088
region_southwest	-0.1289	0.035	-3.680	0.000	-0.198	-0.060
const	7.0308	0.072	97.111	0.000	6.889	7.173

```
=====
Omnibus:                463.941    Durbin-Watson:                2.046
Prob(Omnibus):          0.000    Jarque-Bera (JB):            1674.108
Skew:                   1.679    Prob(JB):                     0.00
Kurtosis:               7.331    Cond. No.                     311.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [66]:

```
# VIF = variance inflation factor= 1/ (1-R^2)
def VIF(df):
    import pandas as pd
    from sklearn.linear_model import LinearRegression

    #initialize dictionaries
    vif_dict, tolerance_dict = {}, {}

    #from input data for each exogenous variable

    for col in df.drop(columns=['const']):
        y= df[col]
        x= df.drop(columns=[col])

        #extract r_squared from the fit

        r_squared = LinearRegression().fit(x, y).score(x, y)

        #calculate VIF
        if r_squared < 1: # Prevent division by zero runtime error
            vif = 1/(1- r_squared)
        else:
            vif = 100
        vif_dict[col] = vif

    #calculate tolerance
```

```
tolerance = 1- r_squared
tolerance_dict[col] = tolerance

    # generate the DataFrame to return
df_output = pd.DataFrame({'VIF': vif_dict, 'Tolerance': tolerance_dict})

return df_output.sort_values(by=['VIF'] , ascending=False)
VIF()
```

```
-----
NameError                                Traceback (most recent call last)
C:\Users\TAWABC~1\AppData\Local\Temp\ipykernel_2528\3845429643.py in <module>
    32
    33     return df_output.sort_values(by=['VIF'] , ascending=False)
--> 34 VIF(X)
```

**NameError:** name 'X' is not defined

In [ ]: